

NFS

Network File System (NFS) is a network file system developed by Sun Microsystems and has the same purpose as SMB. Its purpose is to access file systems over a network as if they were local. However, it uses an entirely different protocol. **NFS** is used between Linux and Unix systems. This means that NFS clients cannot communicate directly with SMB servers. NFS is an Internet standard that governs the procedures in a distributed file system. While NFS protocol version 3.0 (**NFSv3**), which has been in use for many years, authenticates the client computer, this changes with **NFSv4**. Here, as with the Windows SMB protocol, the user must authenticate.

Version	Features
NFSv2	It is older but is supported by many systems and was initially operated entirely over UDP.
NFSv3	It has more features, including variable file size and better error reporting, but is not fully compatible with NFSv2 clients.
NFSv4	It includes Kerberos, works through firewalls and on the Internet, no longer requires portmappers, supports ACLs, applies state-based operations, and provides performance improvements and high security. It is also the first version to have a stateful protocol.

NFS version 4.1 (**RFC 8881**) aims to provide protocol support to leverage cluster server deployments, including the ability to provide scalable parallel access to files distributed across multiple servers (pNFS extension). In addition, NFSv4.1 includes a session trunking mechanism, also known as NFS multipathing. A significant advantage of NFSv4 over its predecessors is that only one UDP or TCP port **2049** is used to run the service, which simplifies the use of the protocol across firewalls.

NFS is based on the **Open Network Computing Remote Procedure Call (ONC-RPC/SUN-RPC)** protocol exposed on **TCP** and **UDP** ports **111**, which uses **External Data Representation (XDR)** for the system-independent exchange of data. The NFS protocol has **no** mechanism for **authentication** or **authorization**. Instead, authentication is completely shifted to the RPC protocol's options. The authorization is taken from the available information of the file system where the server is responsible for translating the user information supplied by the client to that of the file system and converting the corresponding authorization information as correctly as possible into the syntax required by UNIX.

The most common authentication is via UNIX **UID/GID** and **group memberships**, which is why this syntax is most likely to be applied to the NFS protocol. One problem is that the client and server do not necessarily have to have the same mappings of UID/GID to users and groups, and the server does not need to do anything further. No further checks can be made on the part of the server. This is why NFS should only be used with this authentication method in trusted networks.

Default Configuration

NFS is not difficult to configure because there are not as many options as FTP or SMB have. The **/etc/exports** file contains a table of physical filesystems on an NFS server accessible by the clients. The **NFS Exports Table** shows which options it accepts and thus indicates which options are available to us.

Exports File

Exports File

```
dbcyp0n@htb[/htb]$ cat /etc/exports

# /etc/exports: the access control list for filesystems which may be exported
#               to NFS clients.  See exports(5).
#
# Example for NFSv2 and NFSv3:
# /srv/homes      hostname1(rw,sync,no_subtree_check) hostname2(ro,sync,no_subtree_check)
#
# Example for NFSv4:
# /srv/nfs4       gss/krb5i(rw,sync,fsid=0,crossmnt,no_subtree_check)
# /srv/nfs4/homes gss/krb5i(rw,sync,no_subtree_check)
```

The default `exports` file also contains some examples of configuring NFS shares. First, the folder is specified and made available to others, and then the rights they will have on this NFS share are connected to a host or a subnet. Finally, additional options can be added to the hosts or subnets.

Option	Description
<code>rw</code>	Read and write permissions.
<code>ro</code>	Read only permissions.
<code>sync</code>	Synchronous data transfer. (A bit slower)
<code>async</code>	Asynchronous data transfer. (A bit faster)
<code>secure</code>	Ports above 1024 will not be used.
<code>insecure</code>	Ports above 1024 will be used.
<code>no_subtree_check</code>	This option disables the checking of subdirectory trees.
<code>root_squash</code>	Assigns all permissions to files of root UID/GID 0 to the UID/GID of anonymous, which prevents <code>root</code> from accessing files on an NFS mount.

Let us create such an entry for test purposes and play around with the settings.

ExportFS

ExportFS
<pre>root@nfs:~# echo '/mnt/nfs 10.129.14.0/24(sync,no_subtree_check)' >> /etc/exports root@nfs:~# systemctl restart nfs-kernel-server root@nfs:~# exportfs /mnt/nfs 10.129.14.0/24</pre>

We have shared the folder `/mnt/nfs` to the subnet `10.129.14.0/24` with the setting shown above. This means that all hosts on the network will be able to mount this NFS share and inspect the contents of this folder.

Dangerous Settings

However, even with NFS, some settings can be dangerous for the company and its infrastructure. Here are some of them listed:

Option	Description
<code>rw</code>	Read and write permissions.
<code>insecure</code>	Ports above 1024 will be used.

Option	Description
<code>nohide</code>	If another file system was mounted below an exported directory, this directory is exported by its own exports entry.
<code>no_root_squash</code>	All files created by root are kept with the UID/GID 0.

It is highly recommended to create a local VM and experiment with the settings. We will discover methods that will show us how the NFS server is configured. For this, we can create several folders and assign different options to each one. Then we can inspect them and see what settings can have what effect on the NFS share and its permissions and the enumeration process.

We can take a look at the `insecure` option. This is dangerous because users can use ports above 1024. The first 1024 ports can only be used by root. This prevents the fact that no users can use sockets above port 1024 for the NFS service and interact with it.

Footprinting the Service

When footprinting NFS, the TCP ports `111` and `2049` are essential. We can also get information about the NFS service and the host via RPC, as shown below in the example.

Nmap

```

Nmap

dbcyph0n@htb[/htb]$ sudo nmap 10.129.14.128 -p111,2049 -sV -sC

Starting Nmap 7.80 ( https://nmap.org ) at 2021-09-19 17:12 CEST
Nmap scan report for 10.129.14.128
Host is up (0.00018s latency).

PORT      STATE SERVICE VERSION
111/tcp   open  rpcbind 2-4 (RPC #100000)
| rpcinfo:
|   program version    port/proto  service
|   100000   2,3,4        111/tcp     rpcbind
|   100000   2,3,4        111/udp     rpcbind
|   100000   3,4          111/tcp6    rpcbind
|   100000   3,4          111/udp6    rpcbind
|   100003   3            2049/udp    nfs
|   100003   3            2049/udp6   nfs
|   100003   3,4          2049/tcp    nfs
|   100003   3,4          2049/tcp6   nfs
|   100005   1,2,3        41982/udp6  mountd
|   100005   1,2,3        45837/tcp   mountd
|   100005   1,2,3        47217/tcp6  mountd
|   100005   1,2,3        58830/udp   mountd
|   100021   1,3,4        39542/udp   nlockmgr
|   100021   1,3,4        44629/tcp   nlockmgr
|   100021   1,3,4        45273/tcp6  nlockmgr
|   100021   1,3,4        47524/udp6  nlockmgr
|   100227   3            2049/tcp    nfs_acl
|   100227   3            2049/tcp6   nfs_acl
|   100227   3            2049/udp    nfs_acl
|_  100227   3            2049/udp6   nfs_acl
2049/tcp   open  nfs_acl 3 (RPC #100227)
MAC Address: 00:00:00:00:00:00 (VMware)

Service detection performed. Please report any incorrect results at https://nmap.org/submit/ .
Nmap done: 1 IP address (1 host up) scanned in 6.58 seconds

```

The `rpcinfo` NSE script retrieves a list of all currently running RPC services, their names and descriptions, and the ports they use. This lets us check whether the target share is connected to the network on all required ports. Also, for NFS, Nmap has some NSE scripts that can be used for the scans. These can then show us, for example, the `contents` of the share and its `stats`.

Nmap

```
dbcyph0n@htb[/htb]$ sudo nmap --script nfs* 10.129.14.128 -sV -p111,2049
```

```
Starting Nmap 7.80 ( https://nmap.org ) at 2021-09-19 17:37 CEST
Nmap scan report for 10.129.14.128
Host is up (0.00021s latency).
```

```
PORT      STATE SERVICE VERSION
111/tcp   open  rpcbind 2-4 (RPC #100000)
| nfs-ls: Volume /mnt/nfs
|   access: Read Lookup NoModify NoExtend NoDelete NoExecute
| PERMISSION  UID      GID      SIZE  TIME          FILENAME
| rwxrwxrwx   65534   65534   4096   2021-09-19T15:28:17 .
| ?????????? ?      ?      ?      ?      ?      ..
| rw-r--r--   0       0       1872   2021-09-19T15:27:42 id_rsa
| rw-r--r--   0       0       348    2021-09-19T15:28:17 id_rsa.pub
| rw-r--r--   0       0       0      2021-09-19T15:22:30 nfs.share
|_
| nfs-showmount:
|_ /mnt/nfs 10.129.14.0/24
| nfs-statfs:
| Filesystem  1K-blocks  Used      Available  Use%  Maxfilesize  Maxlink
|_ /mnt/nfs    30313412.0  8074868.0  20675664.0  29%   16.0T        32000
| rpcinfo:
| program version  port/proto  service
| 100000  2,3,4          111/tcp    rpcbind
| 100000  2,3,4          111/udp    rpcbind
| 100000  3,4            111/tcp6   rpcbind
| 100000  3,4            111/udp6   rpcbind
| 100003  3              2049/udp   nfs
| 100003  3              2049/udp6  nfs
| 100003  3,4            2049/tcp   nfs
| 100003  3,4            2049/tcp6  nfs
| 100005  1,2,3          41982/udp6 mountd
| 100005  1,2,3          45837/tcp  mountd
| 100005  1,2,3          47217/tcp6 mountd
| 100005  1,2,3          58830/udp  mountd
| 100021  1,3,4          39542/udp  nlockmgr
| 100021  1,3,4          44629/tcp  nlockmgr
| 100021  1,3,4          45273/tcp6 nlockmgr
| 100021  1,3,4          47524/udp6 nlockmgr
| 100227  3              2049/tcp   nfs_acl
| 100227  3              2049/tcp6  nfs_acl
| 100227  3              2049/udp   nfs_acl
|_ 100227  3              2049/udp6  nfs_acl
2049/tcp   open  nfs_acl 3 (RPC #100227)
MAC Address: 00:00:00:00:00:00 (VMware)
```

```
Service detection performed. Please report any incorrect results at https://nmap.org/submit/ .
Nmap done: 1 IP address (1 host up) scanned in 0.45 seconds
```

Once we have discovered such an NFS service, we can mount it on our local machine. For this, we can create a new empty folder to which the NFS share will be mounted. Once mounted, we can navigate it and view the contents just like our local system.

Show Available NFS Shares

Show Available NFS Shares

```
dbcyph0n@htb[/htb]$ showmount -e 10.129.14.128
```

```
Export list for 10.129.14.128:
/mnt/nfs 10.129.14.0/24
```

Mounting NFS Share

Mounting NFS Share

```
dbcyp0n@htb[/htb]$ mkdir target-NFS
dbcyp0n@htb[/htb]$ sudo mount -t nfs 10.129.14.128:/ ./target-NFS/ -o noLOCK
dbcyp0n@htb[/htb]$ cd target-NFS
dbcyp0n@htb[/htb]$ tree .
```

```
.
├── mnt
│   └── nfs
│       ├── id_rsa
│       ├── id_rsa.pub
│       └── nfs.share
```

2 directories, 3 files

There we will have the opportunity to access the rights and the usernames and groups to whom the shown and viewable files belong. Because once we have the usernames, group names, UIDs, and GUIDs, we can create them on our system and adapt them to the NFS share to view and modify the files.

List Contents with Usernames & Group Names

List Contents with Usernames & Group Names

```
dbcyp0n@htb[/htb]$ ls -l mnt/nfs/

total 16
-rw-r--r-- 1 cry0l1t3 cry0l1t3 1872 Sep 25 00:55 cry0l1t3.priv
-rw-r--r-- 1 cry0l1t3 cry0l1t3  348 Sep 25 00:55 cry0l1t3.pub
-rw-r--r-- 1 root      root      1872 Sep 19 17:27 id_rsa
-rw-r--r-- 1 root      root       348 Sep 19 17:28 id_rsa.pub
-rw-r--r-- 1 root      root         0 Sep 19 17:22 nfs.share
```

List Contents with UIDs & GUIDs

List Contents with UIDs & GUIDs

```
dbcyp0n@htb[/htb]$ ls -n mnt/nfs/

total 16
-rw-r--r-- 1 1000 1000 1872 Sep 25 00:55 cry0l1t3.priv
-rw-r--r-- 1 1000 1000  348 Sep 25 00:55 cry0l1t3.pub
-rw-r--r-- 1    0 1000 1221 Sep 19 18:21 backup.sh
-rw-r--r-- 1    0    0 1872 Sep 19 17:27 id_rsa
-rw-r--r-- 1    0    0  348 Sep 19 17:28 id_rsa.pub
-rw-r--r-- 1    0    0     0 Sep 19 17:22 nfs.share
```

It is important to note that if the `root_squash` option is set, we cannot edit the `backup.sh` file even as `root`.

We can also use NFS for further escalation. For example, if we have access to the system via SSH and want to read files from another folder that a specific user can read, we would need to upload a shell to the NFS share that has the `SUID` of that user and then run the shell via the SSH user.

After we have done all the necessary steps and obtained the information we need, we can unmount the NFS share.

Unmounting

Unmounting

```
dbcyp0n@htb[/htb]$ cd ..
dbcyp0n@htb[/htb]$ sudo umount ./target-NFS
```

VPN Servers

 Warning: Each time you "Switch", your connection keys are regenerated and you must re-download your VPN connection file.

All VM instances associated with the old VPN Server will be terminated when switching to a new VPN server.

Existing PwnBox instances will automatically switch to the new VPN server.

us-academy-1

PROTOCOL

☒ UDP 1337 ☐ TCP 443

DOWNLOAD VPN CONNECTION FILE

Start Instance

 / 1 spawns left

Waiting to start..

Questions

Answer the question(s) below to complete this Section and earn cubes!

Target: 10.129.244.90 



Cheat Sheet



Download VPN Connection
File

+ 0 

Enumerate the NFS service and submit the contents of the flag.txt in the "nfs" share as the answer.

Submit your answer here...



Submit

+ 0 

Enumerate the NFS service and submit the contents of the flag.txt in the "nfsshare" share as the answer.

Submit your answer here...

 Cheat Sheet
 Resources
 Go to Questions

Table of Contents

Introduction

Enumeration Principles	✓
Enumeration Methodology	✓

Infrastructure Based Enumeration

Domain Information	✓
Cloud Resources	✓
Staff	✓

Host Based Enumeration

 FTP	✓
 SMB	✓
 NFS	
 DNS	
 SMTP	
 IMAP / POP3	
 SNMP	
 MySQL	
 MSSQL	
 Oracle TNS	
 IPMI	

Remote Management Protocols

Linux Remote Management Protocols	
Windows Remote Management Protocols	

Skills Assessment

 Footprinting Lab - Easy	
 Footprinting Lab - Medium	
 Footprinting Lab - Hard	

My Workstation

OFFLINE

 Start Instance

 / 1 spawns left